

Repaso de algoritmos de machine learning

- 1. SVM
- 2. Decision Tree Classifier
- 3. Random Forest
- 4. K-NN

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
data = pd.read_csv('/content/drug200.csv')
data.head(5)
```

| | Age | Sex | BP | Cholesterol | Na_to_K | Drug |
|---|-----|-----|--------|-------------|---------|-------|
| 0 | 23 | F | HIGH | HIGH | 25.355 | drugY |
| 1 | 47 | M | LOW | HIGH | 13.093 | drugC |
| 2 | 47 | M | LOW | HIGH | 10.114 | drugC |
| 3 | 28 | F | NORMAL | HIGH | 7.798 | drugX |
| 4 | 61 | F | LOW | HIGH | 18.043 | drugY |

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              200 non-null    int64
1   Sex              200 non-null    object
2   BP               200 non-null    object
3   Cholesterol      200 non-null    object
4   Na_to_K          200 non-null    float64
5   Drug             200 non-null    object
dtypes: float64(1), int64(1), object(4)
memory usage: 9.5+ KB
```

```
data.shape
```

(200, 6)

```
data.isnull().sum()
```

```
Age          0
Sex           0
BP            0
Cholesterol  0
Na_to_K       0
Drug          0
dtype: int64
```

Estadística descriptiva

```
data.describe().T
```

| | count | mean | std | min | 25% | 50% | 75% | max |
|---------|-------|-----------|-----------|--------|---------|---------|-------|--------|
| Age | 200.0 | 44.315000 | 16.544315 | 15.000 | 31.0000 | 45.0000 | 58.00 | 74.000 |
| Na_to_K | 200.0 | 16.084485 | 7.223956 | 6.269 | 10.4455 | 13.9365 | 19.38 | 38.247 |

```
#cantidad drogas
data.Drug.value_counts()

drugY    91
drugX    54
```

```

drugA      23
drugC      16
drugB      16
Name: Drug, dtype: int64

```

```

#cantidad bp
data.BP.value_counts()

```

```

HIGH       77
LOW        64
NORMAL     59
Name: BP, dtype: int64

```

```

#cantidad colesterol
data.Cholesterol.value_counts()

```

```

HIGH       103
NORMAL      97
Name: Cholesterol, dtype: int64

```

```

#estudio de skewness
skew_age = data['Age'].skew(axis=0,skipna=True)
skew_na_to_k = data['Na_to_K'].skew(axis=0,skipna=True)

```

```

print(f'skewness age : {skew_age}')
print(f'skewness Na to K : {skew_na_to_k}')

```

```

skewness age : 0.03030835703000607
skewness Na to K : 1.039341186028881

```

```


#graficando
sns.distplot(data.Age)
plt.show()

```

```

/usr/local/lib/python3.8/dist-packages/seaborn/distributions.py:2619: FutureWarning.warn(msg, FutureWarning)

```



```


sns.distplot(data.Na_to_K)
plt.show()

```

```

/usr/local/lib/python3.8/dist-packages/seaborn/distributions.py:2619: FutureWarning.warn(msg, FutureWarning)

```



```

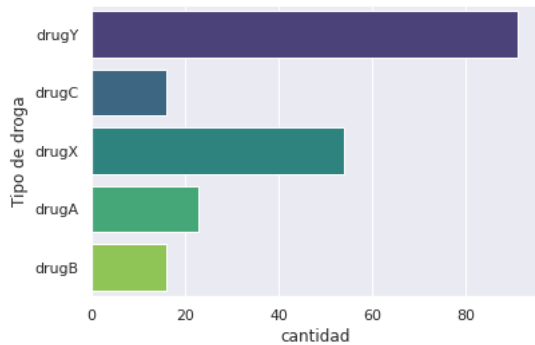
#distribucion por tipo de droga
sns.set_theme(style='darkgrid')

```

```

sns.set_theme(style='darkgrid')
sns.countplot(y='Drug',data=data,palette='viridis')
plt.xlabel('cantidad')
plt.ylabel('Tipo de droga')
plt.show()

```



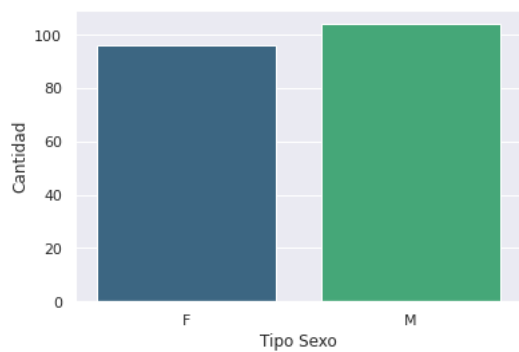
91/200

0.455

```

#Sexo
sns.set_theme(style='darkgrid')
sns.countplot(x='Sex',data=data,palette='viridis')
plt.xlabel('Tipo Sexo')
plt.ylabel('Cantidad')
plt.show()

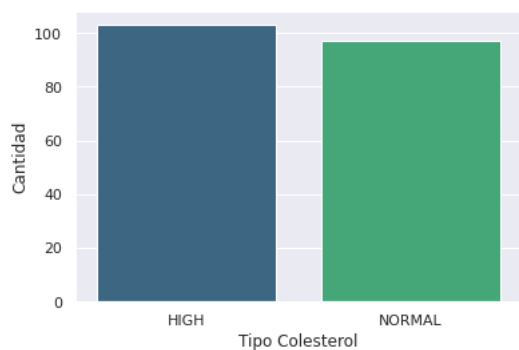
```



```

#colesterol
sns.set_theme(style='darkgrid')
sns.countplot(x='Cholesterol',data=data,palette='viridis')
plt.xlabel('Tipo Colesterol')
plt.ylabel('Cantidad')
plt.show()

```



```

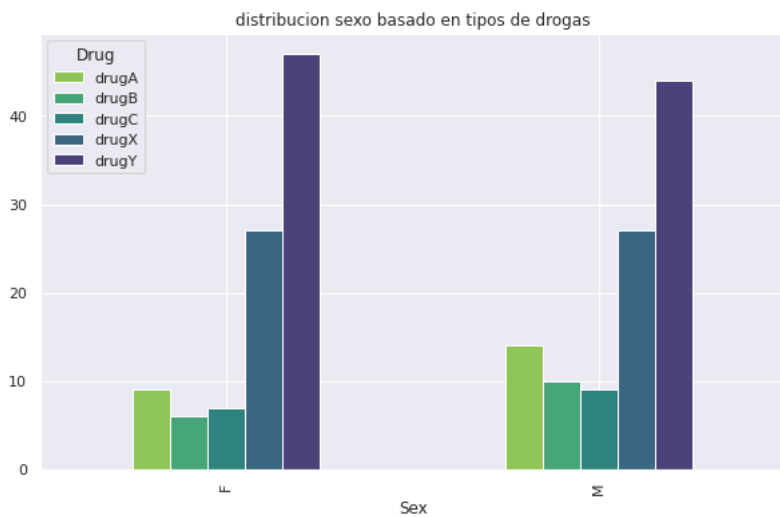
#crosstab pandas

```

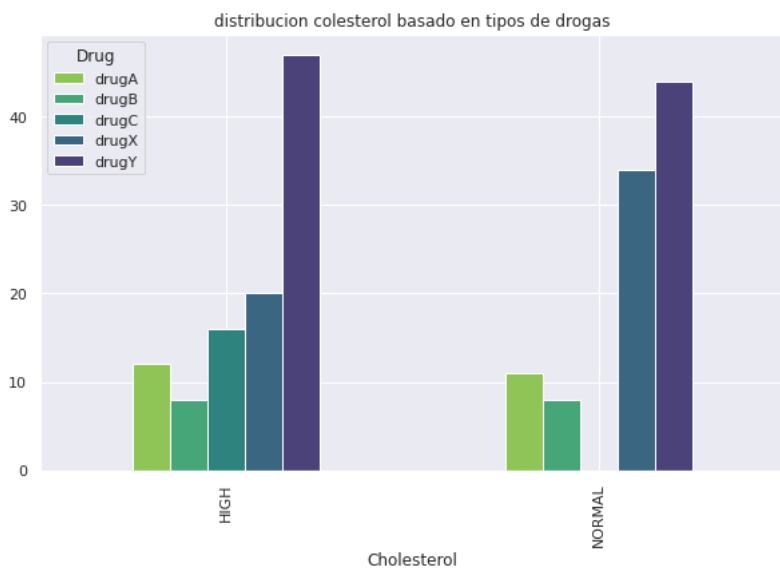
```

pd.crosstab(data.Sex,data.Drug).plot(kind='bar',figsize=(10,6),color=['#8fc456','#47a678','#2e837f','#3b6682','#4b4279'])
plt.title('distribucion sexo basado en tipos de drogas')
plt.show()

```



```
pd.crosstab(data.Cholesterol,data.Drug).plot(kind='bar',figsize=(10,6),color=['#8fc456','#47a678','#2e837f','#3b6682','#4b4279'])
plt.title('distribucion colesterol basado en tipos de drogas')
plt.show()
```



▼ SVM

```
#Creacion de modelos de machine learning
#SVM
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score,classification_report,confusion_matrix
```

```
data.isnull().sum()
```

```
Age      0
Sex      0
BP       0
Cholesterol  0
Na_to_K  0
Drug     0
dtype: int64
```

```
#seleccion de variables
X = data.drop(['Drug'],axis=1)
X = pd.get_dummies(X)
y = data.Drug
```

```
#separacion de data de entrenamiento y testeo
X_train,X_test,y_train,y_test = train_test_split(X,y,random_state=0,test_size=0.5)

SVCModel = SVC(kernel='linear',max_iter=250).fit(X_train,y_train)
y_pred = SVCModel.predict(X_test)


print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| drugA | 0.91 | 1.00 | 0.95 | 10 |
| drugB | 0.86 | 0.86 | 0.86 | 7 |
| drugC | 1.00 | 1.00 | 1.00 | 8 |
| drugX | 1.00 | 1.00 | 1.00 | 28 |
| drugY | 1.00 | 0.98 | 0.99 | 47 |
| accuracy | | | 0.98 | 100 |
| macro avg | 0.95 | 0.97 | 0.96 | 100 |
| weighted avg | 0.98 | 0.98 | 0.98 | 100 |

```
[[10  0  0  0  0]
 [ 1  6  0  0  0]
 [ 0  0  8  0  0]
 [ 0  0  0 28  0]
 [ 0  1  0  0 46]]
/usr/local/lib/python3.8/dist-packages/sklearn/svm/_base.py:284: ConvergenceWarning: Solver terminated early (max_iter=250).
warnings.warn(
```

```
sns.heatmap(pd.DataFrame(confusion_matrix(y_test,y_pred)),annot=True,cmap='Blues')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f44074c8700>



```
print(f'accuracy score : {accuracy_score(y_pred,y_test)}')

accuracy score : 0.98
```

▼ DECISION TREE CLASSIFIER

```
from sklearn.tree import DecisionTreeClassifier,plot_tree

TreeModel = DecisionTreeClassifier(max_leaf_nodes=20).fit(X_train,y_train)
y_pred = TreeModel.predict(X_test)

print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
print(f'accuracy score: {accuracy_score(y_pred,y_test)}')
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| drugA | 0.83 | 1.00 | 0.91 | 10 |
| drugB | 1.00 | 0.71 | 0.83 | 7 |
| drugC | 1.00 | 1.00 | 1.00 | 8 |
| drugX | 1.00 | 0.96 | 0.98 | 28 |
| drugY | 0.98 | 1.00 | 0.99 | 47 |
| accuracy | | | 0.97 | 100 |
| macro avg | 0.96 | 0.94 | 0.94 | 100 |
| weighted avg | 0.97 | 0.97 | 0.97 | 100 |

```
[[10  0  0  0  0]
 [ 2  5  0  0  0]
 [ 0  0  8  0  0]
 [ 0  0  0 27  1]
 [ 0  0  0  0 47]]
accuracy score: 0.97
```

▼ Random Forest

```
from sklearn.ensemble import RandomForestClassifier

RFModel = RandomForestClassifier(max_leaf_nodes=20).fit(X_train,y_train)
y_pred = RFModel.predict(X_test)

print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
print(f'accuracy score: {accuracy_score(y_pred,y_test)}')
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| drugA | 0.83 | 1.00 | 0.91 | 10 |
| drugB | 1.00 | 0.71 | 0.83 | 7 |
| drugC | 1.00 | 1.00 | 1.00 | 8 |
| drugX | 1.00 | 0.96 | 0.98 | 28 |
| drugY | 0.98 | 1.00 | 0.99 | 47 |
| accuracy | | | 0.97 | 100 |
| macro avg | 0.96 | 0.94 | 0.94 | 100 |
| weighted avg | 0.97 | 0.97 | 0.97 | 100 |

```
[[10  0  0  0  0]
 [ 2  5  0  0  0]
 [ 0  0  8  0  0]
 [ 0  0  0 27  1]
 [ 0  0  0  0 47]]
accuracy score: 0.97
```

▼ K-NN

```
from sklearn.neighbors import KNeighborsClassifier

KNNModel = KNeighborsClassifier(n_neighbors=20).fit(X_train,y_train)
y_pred = KNNModel.predict(X_test)

print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
print(f'accuracy score: {accuracy_score(y_pred,y_test)}')
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| drugA | 0.17 | 0.10 | 0.12 | 10 |
| drugB | 0.21 | 0.43 | 0.29 | 7 |
| drugC | 0.00 | 0.00 | 0.00 | 8 |
| drugX | 0.44 | 0.57 | 0.50 | 28 |
| drugY | 0.93 | 0.87 | 0.90 | 47 |
| accuracy | | | 0.61 | 100 |
| macro avg | 0.35 | 0.39 | 0.36 | 100 |
| weighted avg | 0.59 | 0.61 | 0.60 | 100 |

```
[[ 1  0  0  9  0]
 [ 0  3  0  4  0]
 [ 2  0  0  5  1]
 [ 3  7  0 16  2]
 [ 0  4  0  2 41]]
accuracy score: 0.61
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score
_warn_prf(average, modifier, msg_start, len(result))
```



✓ 0 s completado a las 17:07

● ×